

# Intelligent distributed systems

Daniele Fontanelli

Department of Industrial Engineering  
University of Trento

E-mail address: *[daniele.fontanelli@unitn.it](mailto:daniele.fontanelli@unitn.it)*

2022/2023



**UNIVERSITÀ  
DI TRENTO**

**Dipartimento di  
Ingegneria Industriale**

# Outline

- 1 Distributed WLS
- 2 Distributed Kalman Filter

# Outline

1 Distributed WLS

2 Distributed Kalman Filter

# Distributed Least Squares Algorithm

Let us define the same problem we have defined to introduce the LS problem, here reported for clarity:

- We have to determine the value of a *nonrandom* constant parameter;
- We collect the sensor readings from  $m$  *sensors* deployed in an environment by means of a communication system;
- The data are collected in *different time instants*  $t = 0, \dots$  from a subset of  $n(t) \geq m$  of the available sensors. We will refer generically to  $n$  for ease of notation;
- To ensure a correct *sensor fusion*, the data are *timestamped* at the source location and all the sensors are supposed to be *synchronised* at the *Coordinated Universal Time* (UTC), for example using GPS signals;

# Distributed Least Squares Algorithm

- The set of sensor readings is not necessarily the same and, hence, *may change in time*;
- Each sensor provides the measures with a different (possibly time varying) *uncertainty*, i.e., the stochastic process affecting the sensors is *non-stationary*;
- The data are collected by a single entity fusing all the data, i.e., *centralised approach*.

The solution we provided previously is based on the existence of a *central node* that receives all the data and then performs, in a *centralised* way the estimates.

Now, we want to do exactly the same thing in a *distributed way*.

# Distributed Least Squares Algorithm

To properly model this problem, let us define

$$z(i) = H(i)x + \varepsilon(i), \quad i = 1, \dots, n$$

the time varying set of measurements collected at time  $t$ .

Then, let us define the following aggregating vectors

$$Z^n = \begin{bmatrix} z(1) \\ z(2) \\ \vdots \\ z(n) \end{bmatrix}, \quad H^n = \begin{bmatrix} H(1) \\ H(2) \\ \vdots \\ H(n) \end{bmatrix}, \quad \varepsilon^n = \begin{bmatrix} \varepsilon(1) \\ \varepsilon(2) \\ \vdots \\ \varepsilon(n) \end{bmatrix}.$$

# Distributed Least Squares Algorithm

For the covariance matrix of the noises

$$C^n = \begin{bmatrix} C\{\varepsilon(1)\} & 0 & \dots & 0 \\ 0 & C\{\varepsilon(2)\} & \dots & 0 \\ \vdots & \vdots & \ddots & \vdots \\ 0 & 0 & \dots & C\{\varepsilon(n)\} \end{bmatrix}$$

Notice how the noise is assumed uncorrelated in different time instants, while instead it can be correlated for each time instant ( $C\{\varepsilon(i)\}$  is not necessarily diagonal).

# Distributed Least Squares Algorithm

We have then

$$\begin{aligned}
 J(x, n) &= \sum_{i=1}^n (z(i) - H(i)x)^T C \{\varepsilon(i)\}^{-1} (z(i) - H(i)x) = \\
 &= (Z^n - H^n x)^T C^{n-1} (Z^n - H^n x) \Rightarrow \\
 &\Rightarrow \hat{x}^{LS} = \arg \min_x J(x, n) = (H^{nT} C^{n-1} H^n)^{-1} H^{nT} C^{n-1} Z^n.
 \end{aligned}$$

While the covariance matrix of the error is given by

$$P_n = (H^{nT} C^{n-1} H^n)^{-1}$$



# Distributed Least Squares Algorithm

Using the definitions of the matrices given previously, the previous equations can be rewritten as

$$\begin{aligned}\hat{x}^{LS} &= \arg \min_x J(x, n) = (H^{nT} C^{n-1} H^n)^{-1} H^{nT} C^{n-1} Z^n = \\ &= \left( \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} H(i) \right)^{-1} \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} z(i).\end{aligned}$$

Hence, the covariance matrix of the error is given by

$$P_n = \left( \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} H(i) \right)^{-1}$$

# Distributed Least Squares Algorithm

One immediate solution for a *distributed algorithm* to work properly in this case is just using a *flooding algorithm*:

- Each node transmits to all the other nodes in its *communication range* all its data;
- Under mild connectivity properties, sooner or later, *all* the nodes in the network will have *all* the data in the network.
- Now, each node can solve the WLS and hence each node has the *same best LS solution*.

However, such a solution requests a *lot of messages* to travel along the network.

Is there any possibility to reduce this problem?

# Distributed Least Squares Algorithm

Let us start with a *local estimate*.

The sensor considers a local *composite information matrix*  $F_i(k) \in \mathbb{R}^{m \times m}$  and a local *composite information state*  $a_i(k) \in \mathbb{R}^m$ , *one for each node*.

The  $i$ -th sensor makes a measurement and initialises its composite elements as:

$$F_i(0) = H(i)^T C \{\varepsilon(i)\}^{-1} H(i), \text{ and } a_i(0) = H(i)^T C \{\varepsilon(i)\}^{-1} z(i).$$

Therefore, it is able to provide a (first) estimate at time  $k = 0$  using the local batch solution of the WLS, i.e.

$$\hat{x}^{LS}(0) = F_i(0)^{-1} a_i(0).$$

# Distributed Least Squares Algorithm

Instead, the solution of the *centralised* (i.e. *global*) WLS would be given by

$$\begin{aligned}\hat{x}^{LS} &= \left( \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} H(i) \right)^{-1} \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} z(i) = \\ &= \left( \sum_{i=1}^n F_i(0) \right)^{-1} \sum_{i=1}^n a_i(0), \\ P_n &= \left( \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} H(i) \right)^{-1} = \left( \sum_{i=1}^n F_i(0) \right)^{-1}.\end{aligned}$$

The issue is: can we compute the *same* global solution starting from the *local* values of  $F_i(0)$  and  $a_i(0)$ ,  $\forall i$ ?

# Distributed Least Squares Algorithm

Since each sensor has its own

$$F_i(0) = H(i)^T C \{\varepsilon(i)\}^{-1} H(i), \text{ and } a_i(0) = H(i)^T C \{\varepsilon(i)\}^{-1} z(i),$$

it is immediate to notice that

$$\begin{aligned} \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} z(i) &= \sum_{i=1}^n a_i(0) = n \left( \frac{1}{n} \sum_{i=1}^n a_i(0) \right), \\ \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} H(i) &= \sum_{i=1}^n F_i(0) = n \left( \frac{1}{n} \sum_{i=1}^n F_i(0) \right). \end{aligned}$$

Hence, if each node is able to compute the average among  $F_i(0)$  and  $a_i(0)$ ,  $\forall i$ , it can compute the **global** WLS!

This is indeed an **average consensus** problem!

# Distributed Least Squares Algorithm

To account for the design of a protocol ensuring the average consensus, we can make use of the time-varying weighting scheme based on the *maximum-degree weight*

$$q_{ij}(k) = \begin{cases} \frac{1}{n} & \text{if } (j, i) \in \mathcal{E} \text{ and } i \neq j, \\ 1 - \frac{d_i(k)}{n} & \text{if } i = j, \\ 0 & \text{otherwise.} \end{cases}$$

or the *Metropolis weights*

$$q_{ij}(k) = \begin{cases} \frac{1}{\max(d_i(k), d_j(k)) + 1} & \text{if } (j, i) \in \mathcal{E}(k) \text{ and } i \neq j, \\ 1 - \sum_{j=1, i \neq j}^n q_{ij}(k) & \text{if } i = j, \\ 0 & \text{otherwise.} \end{cases}$$

# Distributed Least Squares Algorithm

The update rule is hence based on the *average consensus*

$$\begin{aligned}
 F_i(k+1) &= F_i(k) + \sum_{j=1}^n q_{ij}(k)(F_j(k) - F_i(k)) = \\
 &= q_{ii}(k)F_i(k) + \sum_{j=1, j \neq i}^n q_{ij}(k)F_j(k) \\
 a_i(k+1) &= a_i(k) + \sum_{j=1}^n q_{ij}(k)(a_j(k) - a_i(k)) = \\
 &= q_{ii}(k)a_i(k) + \sum_{j=1, j \neq i}^n q_{ij}(k)a_j(k)
 \end{aligned}$$

# Distributed Least Squares Algorithm

By the previous theorem we have then for  $j$ -th node

$$\lim_{k \rightarrow \infty} F_j(k) = \frac{1}{n} \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} H(i)$$
$$\lim_{k \rightarrow \infty} a_j(k) = \frac{1}{n} \sum_{i=1}^n H(i)^T C \{\varepsilon(i)\}^{-1} z(i)$$



# Distributed Least Squares Algorithm

Hence, for all the  $j = 1, \dots, n$  we have

$$\hat{x}^{LS} = \lim_{k \rightarrow \infty} F_j(k)^{-1} a_j(k).$$

Moreover, at each step (i.e., before reaching the steady state), every node can still compute a *locally unbiased weighted least-square estimate* providing that  $F_i(t)$  is invertible.

As a consequence, *at each iteration* of the consensus protocol, a *new set of measurements* may arrive and fused coherently by means of the WLS. At steady state, all the estimates converge to the *global WLS solution*!

# Distributed Least Squares Algorithm

Dealing with time varying visibility

We now analyse the effect of time-varying graphs for the distributed WLS, introducing a couple of definitions.

## Definition

Let  $\mathcal{G}_i = (\mathcal{N}, \mathcal{E}_i)$ ,  $i = 1, \dots, r$  denote a finite set of graphs sharing *the same set of nodes*. Their *union* is a graph  $\mathcal{G} = \cup_{i=1}^r \mathcal{G}_i = (\mathcal{N}, \cup_{i=1}^r \mathcal{E}_i)$ .

This is obvious by simply computing the overall *t-step transition matrix*.

## Definition

Let  $\mathcal{G}_i = (\mathcal{N}, \mathcal{E}_i)$ ,  $i = 1, \dots, r$  denote a finite set of graphs sharing *the same set of nodes*. The set of graphs  $\{\mathcal{G}_1, \dots, \mathcal{G}_r\}$  is called *jointly connected* if their *union* is *strongly connected*.

# Distributed Least Squares Algorithm

Dealing with time varying visibility

We can now state that the graph is *connected in the long run* if the *infinitely occurring* communication graphs are *jointly connected*.

More precisely, for any graph with *switching topology*, there is a finite set of  $r$  graphs describing its configuration (at most, they are all the possible combination of communication links).

Therefore there is a finite set of associated  $r$  weight matrices  $Q_i$ , determined either with the *maximum degree* or the *Metropolis* algorithms. Then the dynamic becomes

$$x(k+1) = Q_{i(k)}x(k), \text{ with } 1 \leq i(k) \leq r \text{ for all } k.$$

# Distributed Least Squares Algorithm

Dealing with time varying visibility: an example




Figure: A possible sequence of a *jointly connected* graph.

# Distributed Least Squares Algorithm

Dealing with time varying visibility

By the assumption derived previously, the sequence  $\{i(k)\}_{k=0}^{\infty}$  is such that a subset of all the graphs occurs *infinitely often*.

## Theorem

*If the collection of communication graphs that occur infinitely often are jointly connected, then the system*

$$x(k+1) = Q_{i(k)}x(k), \text{ with } 1 \leq i(k) \leq r \text{ for all } k,$$

*converges to*

$$\lim_{k \rightarrow \infty} x(k) = \left( \frac{1}{n} \mathbf{1}^T x(0) \right) \mathbf{1},$$

*or, equivalently*

$$\lim_{k \rightarrow +\infty} \Phi(k) = \frac{1}{n} \mathbf{1} \mathbf{1}^T.$$

# Distributed Least Squares Algorithm

Dealing with time varying visibility

This theorem is given without proof. The proof can be found in:  
Lin Xiao, Stephen Boyd, Sanjay Lall, *A Scheme for Robust Distributed Sensor Fusion Based on Average Consensus*, in Fourth International Symposium on Information Processing in Sensor Networks, pp.63-70, 15 April 2005

The proof is based on the existence of *paracontracting matrices*, i.e.,

$$Mx \neq x \Leftrightarrow \|Mx\| < \|x\|$$

Therefore, even with switching topology and under mild assumptions on communications (e.g., mobile agents that move randomly in the environment with finite communication range), the *optimal estimation problem is solved!*

# Outline

1 Distributed WLS

2 Distributed Kalman Filter

# Distributed Kalman Filter

Let us consider the problem of obtaining a *distributed estimate* of a linear dynamic system.

An example can be the estimate of an agent (e.g. a person) moving in an environment equipped with a surveillance system equipped with a network of  $n$  sensors (e.g. surveillance cameras).

As in the standard Kalman filter notes, the model can then be given by

$$x(k+1) = Ax(k) + Bu(k).$$



# Distributed Kalman Filter

For the distributed Kalman filter, the model turns to

$$\hat{x}(k+1) = A\hat{x}(k) + \nu(k),$$

where  $\nu(k) \sim \mathcal{N}(0, Q)$  models the uncertainties in the motion of the target.

Notice that this is quite obvious if the target is a person, but it is the same if the target is a robot, since it is not possible to get access to the *proprioceptive sensor data*, e.g. odometry.

For the measurement we have then

$$z_i(k) = H_i x_i(k) + \varepsilon_i(k),$$

where  $\varepsilon_i(k) \sim \mathcal{N}(0, R_i)$  models the measurement error.

As customary in the Kalman filter literature, we assume the noises white and uncorrelated.

# Distributed Kalman Filter

For the *centralised Kalman filter* we would have

- *Prediction step:*

$$\hat{x}(k+1)^- = A\hat{x}(k)$$

$$P(k+1)^- = AP(k)A^T + Q$$

- *Update step:*

$$S(k+1) = HP(k+1)^-H^T + R$$

$$W(k+1) = P(k+1)^-H^TS(k+1)^{-1}$$

$$\hat{x}(k+1) = \hat{x}(k+1)^- + W(k+1)(z(k+1) - H\hat{x}(k+1)^-)$$

$$P(k+1) = (I - W(k+1)H)P(k+1)^-$$

where  $R = \text{diag}(R_1, R_2, \dots, R_n)$ .

# Distributed Kalman Filter

The equations can be rewritten more compactly as follows:

- *Prediction step* (the same as before):

$$\hat{x}(k+1)^- = A\hat{x}(k)$$

$$P(k+1)^- = AP(k)A^T + Q$$

- *Update step*:

$$\hat{x}(k+1) = \hat{x}(k+1)^- + P(k+1)^- H^T (HP(k+1)^- H^T + R)^{-1} \cdot (z(k+1) - H\hat{x}(k+1)^-)$$

$$P(k+1) = P(k+1)^- - P(k+1)^- H^T (HP(k+1)^- H^T + R)^{-1} \cdot HP(k+1)^-$$

# Distributed Kalman Filter

Using the *matrix inversion lemma*:

$$(A + BCB^T)^{-1} = A^{-1} - A^{-1}B(B^T A^{-1}B + C^{-1})^{-1}B^T A^{-1},$$

The update equations can be rewritten as follows.

# Distributed Kalman Filter

## State Update

*State update:*

$$\begin{aligned}
 \hat{x}(k+1) &= \hat{x}(k+1)^- + P(k+1)^- H^T (H P(k+1)^- H^T + R)^{-1} \cdot \\
 &\quad \cdot (z(k+1) - H \hat{x}(k+1)^-) = \\
 &= P(k+1) (P(k+1)^- \hat{x}(k+1)^- + H^T R^{-1} z(k+1)) = \\
 &= P(k+1) \left( P(k+1)^- \hat{x}(k+1)^- + \sum_{i=1}^n H_i^T R_i^{-1} z_i(k+1) \right) = \\
 &= P(k+1) (P(k+1)^- \hat{x}(k+1)^- + a(k+1))
 \end{aligned}$$

where

$$a(k+1) = \sum_{i=1}^n H_i^T R_i^{-1} z_i(k+1).$$

# Distributed Kalman Filter

## Covariance Update

*Covariance update:*

$$\begin{aligned}
 P(k+1) &= P(k+1)^- - P(k+1)^- H^T (H P(k+1)^- H^T + R)^{-1} \cdot \\
 &\quad \cdot H P(k+1)^- = \\
 &= (P(k+1)^- + H^T R^{-1} H)^{-1} = \\
 &= \left( P(k+1)^- + \sum_{i=1}^n H_i^T R_i^{-1} H_i \right)^{-1} = \\
 &= (P(k+1)^- + F(k+1))^{-1}
 \end{aligned}$$

where

$$F(k+1) = \sum_{i=1}^n H_i^T R_i^{-1} H_i = F.$$

i.e. it is independent from  $k$  if the number of sensors does not change.

# Distributed Kalman Filter

To summarise:

$$\hat{x}(k+1) = P(k+1) (P(k+1)^- \hat{x}(k+1)^- + a(k+1))$$

$$P(k+1) = (P(k+1)^- + F)^{-1}$$

Notice that the problem is quite similar to the WLS case!

In particular, to correctly execute the Kalman filter, only  $a(k+1)$  and  $F$  involves *message exchanges*, while the other quantities can be computed locally.

The positive aspect is that  $a(k+1)$  and  $F$  can be simply computed using *average consensus*.

# Distributed Kalman Filter

So, as in the WLS case, we start by initialising for the  $i$ -th node:

$$a_i(k+1, 0) = H_i^T R_i^{-1} z_i(k+1) \text{ and } F_i(0) = H_i^T R_i^{-1} H_i,$$

where the 0 stands for the initialisation of the *average consensus* distributed state.

Hence, using the *average linear consensus*, we have

$$\lim_{q \rightarrow \infty} F_i(q) = \frac{1}{n} \sum_{i=1}^n H_i^T R_i^{-1} H_i = \frac{F}{n}$$

$$\lim_{q \rightarrow \infty} a_i(k+1, q) = \frac{1}{n} \sum_{i=1}^n H_i^T R_i^{-1} z_i(k+1) = \frac{a(k+1)}{n}$$

where  $q$  is the number of messages exchanged for the consensus protocol.



# Distributed Kalman Filter

For the correct application of the Kalman Filter we have to assume that *each sensor node knows the number of nodes  $n$*  in order to update the filter coherently, i.e. for the  $i$ -th sensor:

$$\begin{aligned}\hat{x}_i(k+1) &= P_i(k+1) \left( P_i(k+1)^- \hat{x}_i(k+1)^- + n a_i(k+1, q) \right) \\ P_i(k+1) &= \left( P_i(k+1)^- + n F_i(q) \right)^{-1}.\end{aligned}$$

Therefore, before computing the *Update step*, the nodes run a consensus algorithm to reach an agreement for  $a(k+1)$  and  $F$ .

Of course, an agreement is reached after a sufficiently large number of consensus iterations  $q$ , theoretically infinite. If this is not the case, a specific algorithm needs to be designed, such as *finite step convergence* or alternative Kalman approaches.