Intelligent distributed systems

Daniele Fontanelli

Department of Industrial Engineering
University of Trento
E-mail address: *daniele.fontanelli@unitn.it*

2022/2023

UNIVERSITÀ DI TRENTO | Dipartimento di Ingegneria Industriale

Outline

Outline

Unscented Kalman filter

The EKF relies on linearisation to propagate the mean and covariance of the state, which may lead to *overoptimistic estimates*, or even *inconsistency*, when the nonlinearities are severe.
Another approach in this cases is the *Unscented Kalman Filter* (UKF), an extension of the Kalman filter that reduces the linearization errors of the EKF. The use of the UKF can provide significant improvement over the EKF.

Unscented Kalman filter

Propagation of means

Consider a LIDAR acquiring data in *polar coordinates* $p = [r, \theta]^T$, i.e. r being the range and θ the angle.
Assume that both the measurement quantities are affected by the noises η_r and η_θ that are generated by two *wide-sense stationary* and *independent processes* with symmetric pdfs and such that $\mu_{\eta_r} = \mu_{\eta_\theta} = 0$ and variances $\sigma^2_{\eta_r}$ and $\sigma^2_{\eta_\theta}$, i.e. $p_m = [r_m, \theta_m]^T = [r + \eta_r, \theta + \eta_\theta]^T$.
We want to express those quantities in term of the vector of Cartesian coordinates $z = [x, y]^T$, therefore

$$z = \begin{bmatrix} r_m \cos(\theta_m) \\ r_m \sin(\theta_m) \end{bmatrix} = h(p_m)$$

Unscented Kalman filter

Propagation of means

Let us compute a first order approximation:

$$\mu_z = \mathsf{E}\{h(p_m)\} \approx \mathsf{E}\left\{ h(\mathsf{E}\{p_m\}) + \left.\frac{\partial h(p_m)}{\partial p_m}\right|_{\mathsf{E}\{p_m\}} (p_m - \mathsf{E}\{p_m\}) \right\}$$

Since $\mathsf{E}\{p_m\} = p$, we have

$$\mu_z \approx h(p) + \left.\frac{\partial h(p_m)}{\partial p_m}\right|_{p} \mathsf{E}\{p_m - p\} = h(p)$$

Suppose $p = [1, \pi/2]^T$, we have $\mu_z = [0, 1]$.

Unscented Kalman filter

Propagation of means

However, if we consider

$$\mu_x = \mathsf{E}\{x_m\} = \mathsf{E}\{r_m \cos\theta_m\} = \mathsf{E}\{(r+\eta_r)(\cos\theta\cos\eta_\theta - \sin\theta\sin\eta_\theta)\}$$

that, assuming independence of the noises and symmetry of the pdfs and the point $p = [r, \theta]^T = [1, \pi/2]^T$, we have

$$\mu_x = r\cos\theta\mathsf{E}\{\cos\eta_\theta\} = 0$$

If we consider

$$\mu_y = \mathsf{E}\{y_m\} = \mathsf{E}\{r_m \sin\theta_m\} = \mathsf{E}\{(r+\eta_r)(\sin\theta\cos\eta_\theta - \cos\theta\sin\eta_\theta)\}$$

that assuming independence of the noises and symmetry of the pdfs and the point $p = [r, \theta]^T = [1, \pi/2]^T$, we have

$$\mu_y = r\sin\theta\mathsf{E}\{\cos\eta_\theta\} = \mathsf{E}\{\cos\eta_\theta\}$$

Unscented Kalman filter

Propagation of means

Let us assume that η_θ is *uniformly* distributed (if no assumption is made, *it is not possible go any further*), with range in $\pm\eta_\theta^*$.
Therefore, the *analytic solution* is

$$\mu_y = \mathsf{E}\{y_m\} = \mathsf{E}\{\cos\eta_\theta\} = \int_{-\eta_\theta^*}^{\eta_\theta^*} \cos\eta_\theta \frac{1}{2\eta_\theta^*} d\eta_\theta = \frac{\sin\eta_\theta^*}{\eta_\theta^*}$$

Notice that, while with the *first order approximation* we had $\mu_z = [0, 1]$, now instead of 1 we have a number that is equal to 1 *only if* $\eta_\theta^* = 0$, while a value that is < 1, $\forall \eta_\theta^* > 0$.
Hence, the *first order approximation* $\mu_z = [0, 1]$ is *incorrect*!

Unscented Kalman filter

Propagation of means

Let us see what happens if we expand up to the *second order*:

$$\mu_z = \mathsf{E}\left\{h(p_m)\right\} \approx \mathsf{E}\left\{h(\mathsf{E}\left\{p_m\right\}) + \left.\frac{\partial h(p_m)}{\partial p_m}\right|_{\mathsf{E}\{p_m\}} (p_m - \mathsf{E}\left\{p_m\right\})\right\} +$$
$$+ \frac{1}{2}\mathsf{E}\left\{\left.\frac{\partial^2 h(p_m)}{\partial r_m^2}\right|_{\mathsf{E}\{p_m\}} (r_m - \mathsf{E}\left\{r_m\right\})^2\right\} +$$
$$+ \frac{1}{2}\mathsf{E}\left\{\left.\frac{\partial^2 h(p_m)}{\partial \theta_m^2}\right|_{\mathsf{E}\{p_m\}} (\theta_m - \mathsf{E}\left\{\theta_m\right\})^2\right\} +$$
$$+ \mathsf{E}\left\{\left.\frac{\partial^2 h(p_m)}{\partial r_m \partial \theta_m}\right|_{\mathsf{E}\{p_m\}} (r_m - \mathsf{E}\left\{r_m\right\})(\theta_m - \mathsf{E}\left\{\theta_m\right\})\right\}$$

Unscented Kalman filter

Propagation of means

We notice that:

$$\frac{1}{2}\mathsf{E}\left\{\left.\frac{\partial^2 h(p_m)}{\partial r_m^2}\right|_{\mathsf{E}\{p_m\}}(r_m-\mathsf{E}\{r_m\})^2\right\}=\begin{bmatrix}0\\0\end{bmatrix}$$

Then:

$$\mathsf{E}\left\{\left.\frac{\partial^2 h(p_m)}{\partial r_m\partial\theta_m}\right|_{\mathsf{E}\{p_m\}}(r_m-\mathsf{E}\{r_m\})(\theta_m-\mathsf{E}\{\theta_m\})\right\}=\begin{bmatrix}0\\0\end{bmatrix}$$

since the two **rv**s are *uncorrelated*.
Finally:

$$\frac{1}{2}\mathsf{E}\left\{\left.\frac{\partial^2 h(p_m)}{\partial\theta_m^2}\right|_{\mathsf{E}\{p_m\}}(\theta_m-\mathsf{E}\{\theta_m\})^2\right\}=\frac{1}{2}\sigma_{\eta_\theta}^2\begin{bmatrix}0\\-1\end{bmatrix}$$

Unscented Kalman filter

Propagation of means

From this result it follows that the *second order approximation*

$$\mu_z = \begin{bmatrix} 0 \\ 1 \end{bmatrix} + \frac{1}{2}\sigma_{\eta_\theta}^2 \begin{bmatrix} 0 \\ -1 \end{bmatrix}$$

Hence, while $\mu_x = \mathsf{E}\{x_m\} = 0$, we have

$$\mu_y = \mathsf{E}\{y_m\} = \mathsf{E}\{\cos \eta_\theta\} \approx 1 - \frac{\sigma_{\eta_\theta}^2}{2}$$

where the approximation is *up to the second order*.

Unscented Kalman filter

Propagation of means

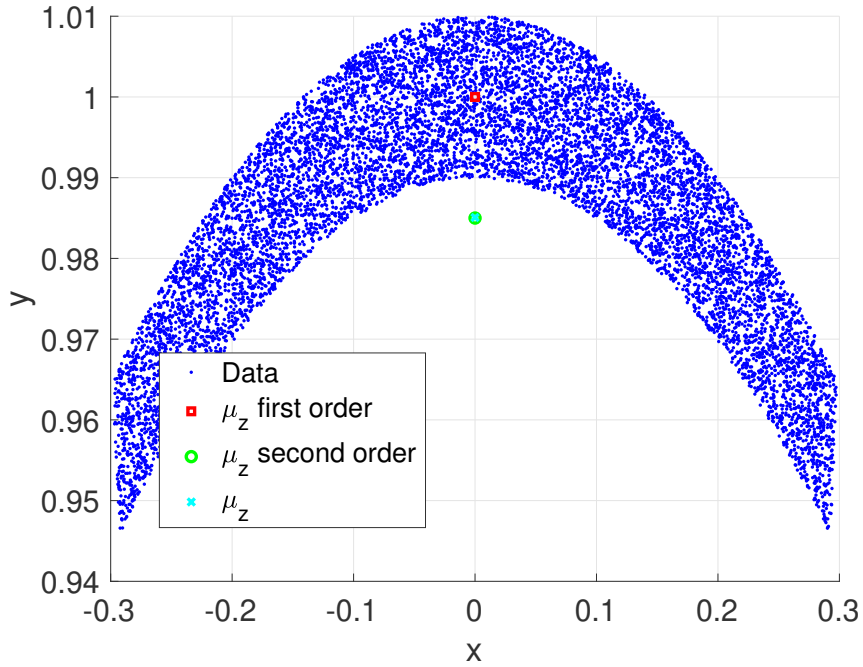

Figure: Approximation of the means for the previous example (with $\eta_r \sim \mathcal{U}(-0.01, 0.01)$ metres and $\eta_\theta \sim \mathcal{U}(-0.4, 0.4)$ radians) for the *first order* Taylor, *second order* Taylor and the solution using a Monte Carlo approach (i.e. *Type B evaluation*).

Unscented Kalman filter

Propagation of covariances

Let us try to compute the covariance of z, i.e.

$$\mathsf{C}\{z\} = \mathsf{E}\left\{(z-\mu_z)(z-\mu_z)^T\right\}$$

If we resort to the *first order approximation*, we have:

$$\mathsf{C}\{z\} \approx H\mathsf{C}\{p_m\}H^T$$

where

$$H = \left.\frac{\partial h(p_m)}{\partial p_m}\right|_{\mathsf{E}\{p_m\}} = \begin{bmatrix} 0 & -1 \\ 1 & 0 \end{bmatrix}$$

Unscented Kalman filter

Propagation of covariances

Notice that

$$\mathsf{C}\left\{p_m\right\} = \mathsf{E}\left\{(p_m - \mathsf{E}\left\{p_m\right\})(p_m - \mathsf{E}\left\{p_m\right\})^T\right\} =$$
$$= \mathsf{E}\left\{\begin{bmatrix}\eta_r \\ \eta_\theta\end{bmatrix}\begin{bmatrix}\eta_r & \eta_\theta\end{bmatrix}\right\} = \begin{bmatrix}\sigma_r^2 & 0 \\ 0 & \sigma_\theta^2\end{bmatrix}$$

which leads to

$$\mathsf{C}\left\{z\right\} \approx H\mathsf{C}\left\{p_m\right\}H^T = \begin{bmatrix}\sigma_\theta^2 & 0 \\ 0 & \sigma_r^2\end{bmatrix}$$

i.e., an *approximated description* of the covariance.

Unscented Kalman filter

Propagation of covariances

Instead, if we make use of the approximation *up to the second order*, we have

$$\mathsf{C}\{z\} = \mathsf{E}\left\{(z-\mu_z)(z-\mu_z)^T\right\} =$$

$$= \mathsf{E}\left\{\begin{bmatrix} r_m\cos\theta_m \\ r_m\sin\theta_m - \left(1-\frac{\sigma_\theta^2}{2}\right)\end{bmatrix}\begin{bmatrix} r_m\cos\theta_m & r_m\sin\theta_m - \left(1-\frac{\sigma_\theta^2}{2}\right)\end{bmatrix}\right\} =$$

$$= \mathsf{E}\left\{\begin{bmatrix} r_m^2\cos\theta_m^2 & r_m^2\cos\theta_m\sin\theta_m - r_m\cos\theta_m\left(1-\frac{\sigma_\theta^2}{2}\right) \\ r_m^2\cos\theta_m\sin\theta_m - r_m\cos\theta_m\left(1-\frac{\sigma_\theta^2}{2}\right) & r_m^2\sin\theta_m^2 + \left(1-\frac{\sigma_\theta^2}{2}\right)^2 - 2r_m\sin\theta_m\left(1-\frac{\sigma_\theta^2}{2}\right)\end{bmatrix}\right\}$$

Unscented Kalman filter

Propagation of covariances

We now notice that if $\eta_r \sim \mathcal{U}(-r^*, r^*)$ with variance σ_r^2, $\eta_\theta \sim \mathcal{U}(-\eta_\theta^*, \eta_\theta^*)$ with variance σ_θ^2 and $p = [r, \theta]^T = [1, \pi/2]$, we have

$$\mathsf{E}\left\{r_m^2 \cos\theta_m^2\right\} = \mathsf{E}\left\{r_m^2\right\} \mathsf{E}\left\{\cos\theta_m^2\right\}$$

$$\mathsf{E}\left\{r_m^2\right\} = \mathsf{E}\left\{r^2\right\} + \mathsf{E}\left\{\eta_r^2\right\} + 2r\mathsf{E}\left\{\eta_r\right\} = 1 + \sigma_r^2$$

$$\mathsf{E}\left\{\cos\theta_m^2\right\} = \frac{1}{2}(1 + \mathsf{E}\left\{\cos 2\theta_m\right\})$$

$$\mathsf{E}\left\{\sin\theta_m^2\right\} = \frac{1}{2}(1 - \mathsf{E}\left\{\cos 2\theta_m\right\})$$

$$\mathsf{E}\left\{\cos 2\theta_m\right\} = -\mathsf{E}\left\{\cos 2\eta_\theta\right\}$$

$$\mathsf{E}\left\{\cos 2\eta_\theta\right\} \approx 1 - 2\sigma_\theta^2$$

$$\mathsf{E}\left\{\cos\theta_m\right\} = 0$$

$$\mathsf{E}\left\{\sin\theta_m\right\} = \mathsf{E}\left\{\cos\eta_\theta\right\} \approx 1 - \frac{\sigma_\theta^2}{2}$$

$$\mathsf{E}\left\{\cos\theta_m \sin\theta_m\right\} = -\sin_\theta^2 \mathsf{E}\left\{\cos\eta_\theta \sin\eta_\theta\right\} = 0$$

Unscented Kalman filter

Propagation of covariances

Substituting

$$\begin{aligned}\mathsf{C}\{z\} &= \mathsf{E}\left\{(z-\mu_z)(z-\mu_z)^T\right\} = \\ &= \begin{bmatrix}(1+\sigma_r^2)\sigma_\theta^2 & 0 \\ 0 & (1+\sigma_r^2)(1+\sigma_\theta^2) - \left(1-\frac{\sigma_\theta^2}{2}\right)^2\end{bmatrix}\end{aligned}$$

Unscented Kalman filter

Propagation of covariances

Instead, if we make use of the *analytic solution* with no approximation (which *is only valid* for uniform pdfs), we have

$$\mathsf{C}\{z\} = \mathsf{E}\left\{(z-\mu_z)(z-\mu_z)^T\right\} =$$

$$= \mathsf{E}\left\{\begin{bmatrix} r_m\cos\theta_m \\ r_m\sin\theta_m - \frac{\sin\eta_\theta^*}{\eta_\theta^*}\end{bmatrix}\begin{bmatrix} r_m\cos\theta_m & r_m\sin\theta_m - \frac{\sin\eta_\theta^*}{\eta_\theta^*}\end{bmatrix}\right\} =$$

$$= \mathsf{E}\left\{\begin{bmatrix} r_m^2\cos\theta_m^2 & r_m^2\cos\theta_m\sin\theta_m - r_m\cos\theta_m\frac{\sin}{\eta} \\ r_m^2\cos\theta_m\sin\theta_m - r_m\cos\theta_m\frac{\sin\eta_\theta^*}{\eta_\theta^*} & r_m^2\sin\theta_m^2 + \frac{\sin\eta_\theta^{*2}}{\eta_\theta^{*2}} - 2r_m\sin\theta_m\frac{s}{}\end{bmatrix}\right.$$

Unscented Kalman filter

Propagation of covariances

With respect to the previous case we only have the following differences

$$\mathsf{E}\left\{\cos 2\eta_\theta\right\} = \frac{\sin 2\eta_\theta^*}{2\eta_\theta^*}$$

$$\mathsf{E}\left\{\sin\theta_m\right\} = \mathsf{E}\left\{\cos\eta_\theta\right\} = \frac{\sin\eta_\theta^*}{\eta_\theta^*}$$

Unscented Kalman filter

Propagation of covariances

Hence, substituting we have

$$\mathsf{C}\{z\} = \mathsf{E}\left\{(z-\mu_z)(z-\mu_z)^T\right\} =$$

$$= \begin{bmatrix} \frac{1}{2}(1+\sigma_r^2)(1-\frac{\sin 2\eta_\theta^*}{2\eta_\theta^*}) & 0 \\ 0 & \frac{1}{2}(1+\sigma_r^2)(1+\frac{\sin 2\eta_\theta^*}{2\eta_\theta^*}) - \frac{\sin {\eta_\theta^*}^2}{{\eta_\theta^*}^2} \end{bmatrix}$$

Unscented Kalman filter

Propagation of covariances

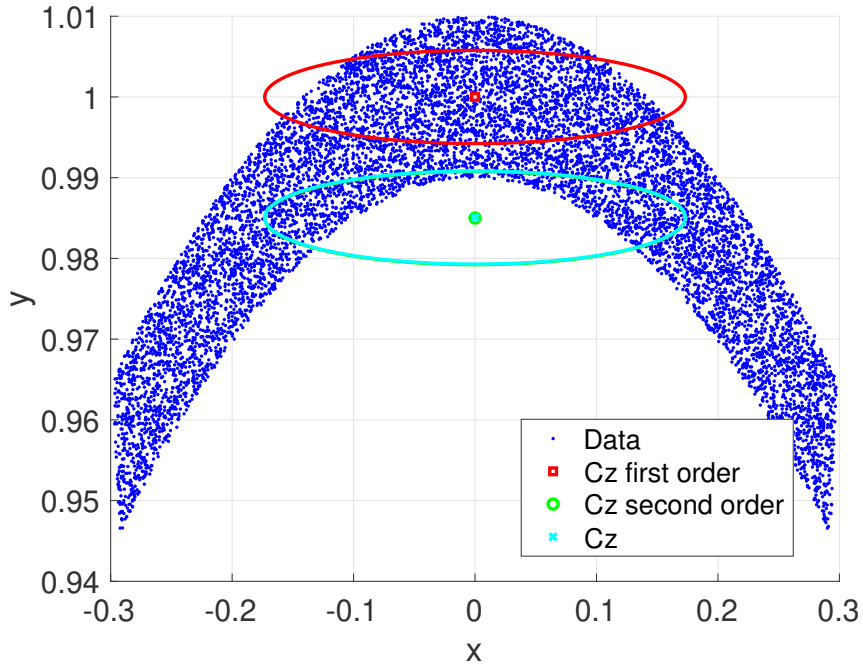

Figure: Approximation of the covariances for the previous example (with $\eta_r \sim \mathcal{U}(-0.01, 0.01)$ metres and $\eta_\theta \sim \mathcal{U}(-0.4, 0.4)$ radians) for the *first order* Taylor, *second order* Taylor and the *analytic* solution using a Monte Carlo approach (i.e. *Type B evaluation*).

Unscented Kalman filter

The main message is that a second order approximation *fits better* when the nonlinearities becomes relevant.
The *mean* and the *variance* are not sufficient to fully describe the pdf in many cases.

Unscented Kalman filter

Propagation of covariances

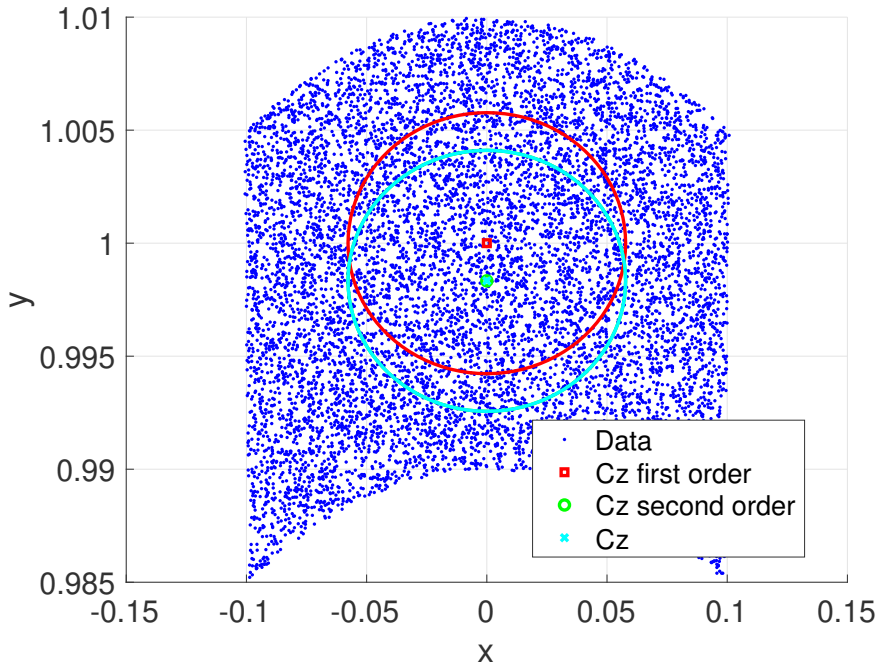

Figure: Approximation of the covariances for the previous example when $\eta_\theta \sim \mathcal{U}(-0.1, 0.1)$ radians.

Unscented Kalman filter

Unscented Transformation

It is in general quite difficult to determine the pdf $p(y)$ of y if $y = f(x)$. However:

- It is easy to compute $y^{(i)} = f(x^{(i)})$, where $x^{(i)}$ is a point extracted from $p(x)$;
- If we collect a set of $2n$ points $x^{(i)}$ (where n is the dimension of x), generating $y^{(i)} = f(x^{(i)})$, $i = 1, \ldots, 2n$, such that from $y^{(i)}$ it is possible to *approximate* $p(y)$, then the problem turns out to be feasible;
- Such points $x^{(i)}$ are called *sigma-points*.

Unscented Kalman filter

Unscented Transformation

More precisely:

- Suppose we know $\mathsf{E}\{x\} = \mu_x$ and $\mathsf{C}\{x\} = P_x$;
- We derive the $2n$ *sigma-points* $x^{(i)}$ such that the their *sample mean* and *sample covariance matrix* are μ_x and P_x;
- We then derive $y^{(i)} = f(x^{(i)})$, $i = 1, \ldots, 2n$;
- The *sample mean* $\hat{\mu}_y$ and *sample covariance matrix* $\hat{P}_y$ will then be, with a reasonable approximation, the actual μ_y and P_y of $p(y)$.

The described algorithm is called the *unscented transformation*.

Unscented Kalman filter

Unscented Transformation

One way to generate the $2n$ *sigma-points* is the following:

- We denote with $\sqrt{nP_x}$ the square root of nP_x such that $\sqrt{nP_x}^T \sqrt{nP_x} = nP_x$ (this can be obtained with the *Cholesky decomposition*);
- We denote with $\left(\sqrt{nP_x}\right)_i$ the i-th row of $\sqrt{nP_x}$;
- The *sigma-points* are then possibly given by

$$
\begin{aligned}
x^{(i)} &= \mu_x + \tilde{x}^{(i)}, \quad i = 1, \dots, 2n \\
\tilde{x}^{(i)} &= \left(\sqrt{nP_x}\right)_i^T, \quad i = 1, \dots, n \\
\tilde{x}^{(i+n)} &= -\left(\sqrt{nP_x}\right)_i^T, \quad i = 1, \dots, n
\end{aligned}
$$

Unscented Kalman filter

Unscented Transformation

Given the *sigma-points*, is then possible to estimate the mean $\hat{\mu}_y$ and covariance matrix $\hat{P}_y$ of the actual μ_y and P_y as follows:

- $\hat{\mu}_y = \frac{1}{2n}\sum_{i=1}^{2n} y^{(i)} = \frac{1}{2n}\sum_{i=1}^{2n} f(x^{(i)})$;
- $\hat{P}_y = \frac{1}{2n}\sum_{i=1}^{2n}(y^{(i)} - \hat{\mu}_y)(y^{(i)} - \hat{\mu}_y)^T = \frac{1}{2n}\sum_{i=1}^{2n}(f(x^{(i)}) - \hat{\mu}_y)(f(x^{(i)}) - \hat{\mu}_y)^T$.

Unscented Kalman filter

Unscented Transformation

The main advantages of the *unscented transformation* are the following:

- There is no need to linearize the functions (as it happens, e.g., in the EKF);
- The approximation given by $\hat{\mu}_y$ is correct up to the *second order* of the Taylor expansion (not the first, as in the classic linearisation). To prove this, it is sufficient to compute the Taylor expansion of $\hat{\mu}_y$ around μ_x and then noticing that by the peculiar choice of the *sigma-points* (i.e., $\tilde{x}^{(i)} = -\tilde{x}^{(i+n)}$), the odd terms of the expansion cancels out;
- Instead, for the approximation given by $\hat{P}_y$, the order is the same as for the linearised version. Nonetheless, there is some level of knowledge of higher order terms (e.g., at least the sign is estimated correctly), hence a higher degree of approxiamtion is obtained.

Unscented Kalman filter

The definition of the UKF then follows directly from the application of the *unscented transformation* to a Kalman Filter to propagate both the model transformation and the measurement transformation.
Therefore, given the following system definition

$$\begin{cases} x_{k+1} = f_k(x_k, u_k) + \nu_k \\ z_k = h_k(x_k) + \varepsilon_k \end{cases}$$

where ν_k and ε_k are two zero-mean uncertainties generated by a white stochastic process, with covariance matrices Q_k and R_k, respectively. Notice that the system can be time-varying, but the noises are at the moment considered as *additive* (e.g., a system that is affine in the *uncertainties*).

Unscented Kalman filter

Initialisation step:

- $$\hat{x}_0 = \mathsf{E}\{x_0\} \text{ and } P_0 = \mathsf{E}\left\{(x_0 - \hat{x}_0)(x_0 - \hat{x}_0)^T\right\}.$$

Unscented Kalman filter

Prediction step: Given $\hat{x}_{k-1}$ and P_{k-1}, compute

- *Sigma-points*:

$$\hat{x}_{k-1}^{(i)} = \hat{x}_{k-1} + \tilde{x}^{(i)}, \quad i = 1, \dots, 2n$$
$$\tilde{x}^{(i)} = \left(\sqrt{nP_{k-1}}\right)_i^T, \quad i = 1, \dots, n$$
$$\tilde{x}^{(i+n)} = -\left(\sqrt{nP_{k-1}}\right)_i^T, \quad i = 1, \dots, n$$

- Transform the *sigma-points* according to the system dynamics, i.e.

$$\hat{x}_k^{(i)} = f_k(\hat{x}_{k-1}^{(i)}, u_k).$$

Unscented Kalman filter

Prediction step:

- Compute the *predicted* estimates

$$\hat{x}_k^- = \frac{1}{2n} \sum_{i=1}^{2n} \hat{x}_k^{(i)}.$$

- Estimate the *predicted* covariance matrix of the estimation error

$$\hat{P}_k^- = \frac{1}{2n} \sum_{i=1}^{2n} (\hat{x}_k^{(i)} - \hat{x}_k^-)(\hat{x}_k^{(i)} - \hat{x}_k^-)^T + Q_{k-1}.$$

Unscented Kalman filter

Measurement update step:

- *Sigma-points*:

$$\hat{x}_k^{(i)} = \hat{x}_k^- + \tilde{x}^{(i)}, \quad i = 1, \dots, 2n$$
$$\tilde{x}^{(i)} = \left(\sqrt{nP_k^-}\right)_i^T, \quad i = 1, \dots, n$$
$$\tilde{x}^{(i+n)} = -\left(\sqrt{nP_k^-}\right)_i^T, \quad i = 1, \dots, n$$

  This step can use the same sigma-points of the prediction to save computation time, *at the price of a reduced performance*.

- Transform the *sigma-points* according to the system measurement function, i.e.

$$\hat{z}_k^{(i)} = h_k(\hat{x}_k^{(i)}).$$

Unscented Kalman filter

Measurement update step:

- Compute the *output* mean

$$\hat{z}_k = \frac{1}{2n} \sum_{i=1}^{2n} \hat{z}_k^{(i)}.$$

- Estimate the *innovation* covariance matrix

$$\hat{S}_z = \frac{1}{2n} \sum_{i=1}^{2n} (\hat{z}_k^{(i)} - \hat{z}_k)(\hat{z}_k^{(i)} - \hat{z}_k)^T + R_k.$$

- Estimate the *innovation* and state cross-covariance matrix

$$\hat{P}_{x,z} = \frac{1}{2n} \sum_{i=1}^{2n} (\hat{x}_k^{(i)} - \hat{x}_k^-)(\hat{z}_k^{(i)} - \hat{z}_k)^T.$$

Unscented Kalman filter

Measurement update step:

- Compute the UKF gain

$$W_k = \hat{P}_{x,z}\hat{S}_z^{-1}.$$

- Update the estimates

$$\hat{x}_k = \hat{x}_k^- + W_k(z_k - \hat{z}_k).$$

- Update the estimation error covariance matrix

$$P_k = P_k^- - W_k\hat{S}_z W_k^T.$$

- Then, starts over.

Unscented Kalman filter

Remark

If the uncertainties are not additive, i.e.

$$\begin{cases} x_{k+1} = f_k(x_k, u_k, \nu_k) \\ z_k = h_k(x_k, \varepsilon_k) \end{cases}$$

the problem can be solved by adding the uncertainties in the augmented state, i.e.

$$q_k = \begin{bmatrix} x_k \\ \nu_k \\ \varepsilon_k \end{bmatrix}.$$

Unscented Kalman filter

Remark

In such a case, we have

$$\hat{q}_0 = \begin{bmatrix} E\{x_0\} \\ 0 \\ 0 \end{bmatrix}$$

and

$$P_0 = \begin{bmatrix} E\{(x_0 - E\{x_0\})(x_0 - E\{x_0\})^T\} & 0 & 0 \\ 0 & Q_0 & 0 \\ 0 & 0 & R_0 \end{bmatrix}.$$

*Details can be found in S. Julier and J. Uhlmann, "*Unscented filtering and nonlinear estimation*", Proceedings of the IEEE, 92(3), pp. 401-422 (March 2004).*

Unscented Kalman filter

Remark

There are other possible unscented transformation that can be applied. The most popular are the generalised unscented transformation (using $2n+1$ points, an additional tuning parameter κ and generic weights of the sigma-points instead of $1/2n$), the simplex unscented transformation (which minimises the number of sigma-points to $n+1$) and the spherical unscented transformation (developed with the goal of rearranging the sigma-points of the simplex algorithm in order to obtain better numerical stability).

Unscented Kalman filter

Remark

UKF can be "easily" distributed among several robots using similar steps as for the EKF.
The interested reader may have a look at:
Wenling Li and Yingmin Jia, Consensus-Based Distributed Multiple Model UKF for Jump Markov Nonlinear Systems, *IEEE Transactions on Automatic Control, Vol. 57, No. 1, January 2012.*

Outline

Particle filter

Particle filters are a kind of *brute force* approaches developed in 40s, but actually applied from the 80s due to their relatively high computational burden.
The idea behind the *particle filter* is the same of the *frequentist* definition of probabilities: when a problem becomes too complicated, with many possible outcomes, it is impossible to give an analytic description, however an estimate of the probability may be given by *observation*.
We saw that UKF can be used when the EKF is a too harsh approximation. Nevertheless, both UKF and EKF need tuning. When the nonlinearities become too severe or the system is hard to be modelled, that is when the *Particle filter* comes into play.

Particle filter

We recall here the recursive Bayesian filter steps:

- The first computes the *new* estimate given all the previous measurements. In doing so, it makes use of the *system dynamics*, i.e.

$$p(x(k)|Z^{k-1}) = \int_{-\infty}^{+\infty} p(x(k)|x(k-1))p(x(k-1)|Z^{k-1})dx(k-1).$$

  This step is called *prediction*: it predicts the next value $x(k)$ given the previous knowledge $x(k-1)$;
- The second refines the prior with the *new set of measurements* (i.e. the new gained knowledge) using the likelihood function, i.e.

$$p(x(k)|Z^k) = \frac{p(z(k)|x(k))p(x(k)|Z^{k-1})}{p(z(k)|Z^{k-1})}.$$

  Since this second step updates the prediction with new knowledge, it is called *update*.

Particle filter

For the denominator of the Bayes equation (i.e., the normalisation factor), we consider the *Total Probability Law* and the Markovian property, thus having

$$p(z(k)|Z^{k-1}) = \int_{-\infty}^{+\infty} p(z(k)|x(k))p(x(k)|Z^{k-1})dx(k),$$

that is, as usual, the integral of the numerator (normalisation).

Particle filter

Linear system with Gaussian uncertainties

From the definition of the dynamic systems

$$x(k+1) = Ax(k) + \nu(k),$$

we have immediately that

$$\mathsf{V}\left\{x(k+1)\right\} = AP_kA^T + Q,$$

hence $\Sigma = Q$ i.e., the covariance matrix of the r.v. $x(k+1)$ given the model!
Notice how the uncertainty inevitably increases in the *prediction* step since AP_kA^T and Q are both p.d.

Particle filter

Linear system with Gaussian uncertainties

To summarise, assuming that $\mathsf{E}\{\nu(k)\nu(j)\} = Q(k)\delta_{k,j}$, $\nu(k) \sim \mathcal{N}(0, Q(k))$, $\mathsf{E}\{\varepsilon(k)\varepsilon(j)\} = R(k)\delta_{k,j}$, $\varepsilon(k) \sim \mathcal{N}(0, R(k))$, $\mathsf{E}\{\nu(k)\varepsilon(j)\} = 0$, $\forall i, j$, and

$$\begin{aligned} x(k+1) &= A(k)x(k) + B(k)u(k) + G(k)\nu(k), \\ z(k) &= H(k)x(k) + F(k)\varepsilon(k), \end{aligned}$$

we have

$$\begin{aligned} p(x(k)|Z^k) &= \mathcal{N}(\hat{x}(k), P(k)), \\ p(x(k+1)|x(k)) &= \mathcal{N}(A(k)\hat{x}(k) + B(k)u(k), G(k)Q(k)G(k)^T), \\ p(z(k+1)|x(k+1)) &= \mathcal{N}(H(k+1)\hat{x}(k+1), R(k+1)), \\ p(x(k+1)|Z^{k+1}) &= \mathcal{N}(\hat{x}(k+1), P(k+1)). \end{aligned}$$

Particle filter

In practice the *particle filter* represents a numerical implementation of a *Bayesian estimator*, considering a generic nonlinear system:

$$\begin{cases} x_{k+1} = f_k(x_k, u_k, \nu_k) \\ z_k = h_k(x_k, \varepsilon_k) \end{cases}$$

where ν_k and ε_k are assumed to be *independent* and *white*.
The main objective of the *particle filter* is to approximate numerically $p(x(k+1)|Z^{k+1})$ using all the measurements up to time $k+1$ in a recursive way.
Notice that we assume that the first measurement is available at time $k = 1$, hence the filter is initialised trivially assuming $p(x_0) = p(x_0|Z^0)$, i.e., a prior.

Particle filter

Moreover, we assume that the pdfs involved in the *particle filter* estimator are generic.
Hence, they can be *multimodal*!
Notice that in this case, assuming as the *best* estimate the *mean* can be a very poor choice (e.g., image a multimodal pdf symmetric with respect to the mean).

Particle filter

Initialisation step:

- At the beginning, a given number of n state vectors (i.e., *particles*) $x_0^{(i)}$, $i = 1, \ldots, n$, are randomly generated from $p(x_0)$, which is *assumed to be known*.

Particle filter

Prediction step:

- The *particles* are then propagated following the system dynamics, i.e.

$$x_{k+1}^{(i),-} = f_k(x_k^{(i)}, u_k, \nu_k),$$

where ν_k is *randomly generated* form the *known* pdf $p(\nu_k)$.

Particle filter

Measurement step: given $z_k = h_k(x_k, \varepsilon_k)$

- When the measurement z_{k+1} is available, the *conditional relative likelihood* is computed for each particle $x_{k+1}^{(i),-}$, that is we *numerically* evaluate $p(z_{k+1}|x_{k+1}^{(i),-})$.
- In practice, given the sensor readings $z^\star$, we compute numerically

$$\lambda_i = \Pr\left[z_{k+1} = z^\star | x_{k+1} = x_{k+1}^{(i),-}\right] \approx \Pr\left[\varepsilon_k = z^\star - h_k(x_{k+1}^{(i),-})\right],$$

that is the *probability that the i-th particle is correct given z_{k+1}*.

Particle filter

Measurement step:

- The probabilities are then normalised, i.e.

$$\lambda_i = \frac{\lambda_i}{\sum_{i=1}^{n} \lambda_i}.$$

- The particles are then *resampled* using λ_i. This resampling can be done in several ways. For example, for the i-th sample:
  - Generate a random number $y \sim \mathcal{U}(0,1)$. If $\sum_{i=1}^{j-1} \lambda_i < y$, but $\sum_{i=1}^{j} \lambda_i \geq y$, then $x_{k+1}^{(i)} = x_{k+1}^{(j),-}$ with probability π_j.
  - This resampling can be shown to converge towards the actual $p(x_{k+1}|Z^{k+1})$ when $n \to +\infty$.

Particle filter

Measurement step:

- After the resampling, it is then possible to have an estimate, e.g. the *conditional mean* $\mathrm{E}\left\{x_{k+1}|Z^{k+1}\right\}$, with the caveat mentioned previously for multimodal pdfs.

Particle filter

Example (Particle relative likelihood)

Consider $z_{k+1} = h_k(x_k) + \varepsilon_k$, where $z_{k+1} \in \mathbb{R}^m$, $\varepsilon_k \sim \mathcal{N}(0, R_k)$ and $z^\star$ is the actual sensor reading.
We have immediately that:

$$\lambda_i \propto \frac{1}{\sqrt{|2\pi R_k|}} e^{-\frac{1}{2}(z^\star - h(x_{k+1}^{(j),-}))^T R_k^{-1} (z^\star - h(x_{k+1}^{(j),-}))}.$$

Notice that on the righthand side we do not have a probability (we missed the integral), but it is proportional to the probability.
This is the reason why we are normalising the probabilities.

Particle filter

Some additional comments:

- The *particle filter* suffers of *sample impoverishment*.
  - This is induced by the fact that the predicted pdf $p(x_{k+1}^{(j),-}|Z^k)$ is resampled over the relative likelihood $p(z_{k+1}|x_{k+1}^{(i),-})$, so it is expected that a large faction of the apriori particles will have low likelihood probability. This will happen sooner or later.
  - The problem becomes even harsher if there are *modelling erros* in the output models.
  - Solution: *roughening*, *prior editing*, *regularisation*, *Markov Chain Monte Carlo resampling*, *auxiliary particle filtering*.

Particle filter

- Some solutions exist combining particle filtering with EKF or UKF, i.e.:
  - Each particle is *updated at the measurement time* using the EKF or the UKF;
  - Resampling is then performed using the measurement;
  - This is like running multiple KF, EKF or UKF Kalman filters (one for each particle) and then adding a resampling step after each measurement.

Particle filter

Remark

Particle filtering goes by many names: sequential importance sampling, bootstrap filtering, condensation algorithm, interacting particle approximations, Monte Carlo filtering, sequential Monte Carlo filtering, survival of the fittest.

Remark

If the system is generic, EKF works fine, UKF even better and particle filters even better. The increasing performance is paid with an increasing computation time.
If the system is linear and Gaussian, KF is the best (EKF is the same of the KF), while UKF and particle filters works as the KF but with an increasing computational load.

Particle filter

Remark

Particle filtering can be hardly distributed on several robots, due to the large amount of computations, hence communications, needed for the executions on several robots.
Nonetheless, once the estimate is available in one robot, it is just a matter of computing an, e.g., WLS among the estimates.